

# Data Mining & Machine Learning

---

CS37300

Purdue University

September 25, 2017

# Searching over models/patterns

---

- Consider a **space** of possible models  $M=\{M_1, M_2, \dots, M_k\}$  with parameters  $\theta$
- Search could be over model structures or parameters, e.g.:
  - **Parameters:** In a linear regression model, what are regression coefficients ( $\beta$ ) that minimize squared loss on the training data?
  - **Model structure:** In a decision trees, what is the tree structure that minimizes 0/1 loss on the training data?

# Combinatorial optimization

# Optimization

---

- **Non-smooth** functions:
  - If the function is *discrete*, then traditional optimization methods that rely on smoothness are not applicable (e.g., gradient descent needs the derivative). Instead we need to use **combinatorial optimization**
  - *Example: Choosing what features (structure) to add to a decision tree*

# Search algorithms for discrete spaces

---

- Conduct the search by:
  - Considering a particular state (*model*)
  - Testing to see if it is the goal state (*model with maximum score*)
  - And if not, expand the current state to generate successor states by applying all possible actions  
(*determine alternative models to consider next*)
- Search strategies differ in their choice of how to expand states

# Heuristic search

---

- Typically, there is an exponential number of models in the (discrete) search space, making it intractable to exhaustively search the space
  - Thus, it is generally impossible to return a model that is **guaranteed** to have the best score
- Instead, we have to resort to **heuristic** search techniques
  - Methods are evaluated experimentally and shown to have good performance on average
  - **Greedy** search: Given a current model  $M$ , look for other models near  $M$  and move to the best of these (if any have a score better than  $M$ )

# Greedy search

---

- Choose an initial state  $M^0$  corresponding to a particular model structure (e.g., an empty tree)
- Let  $M^i$  be the model structure location at the  $i$ -th iteration
- For each iteration  $i$ 
  - Construct all possible models  $\{M^{j1}, \dots, M^{jk}\}$  adjacent to  $M^i$  (as defined by search operators)
  - Evaluate scores for all models  $\{M^{j1}, \dots, M^{jk}\}$
  - Choose to move to the adjacent model with best score:  $M^{i+1} = M^{j.\text{best}}$
  - Repeat until there is no possible further improvement in the score

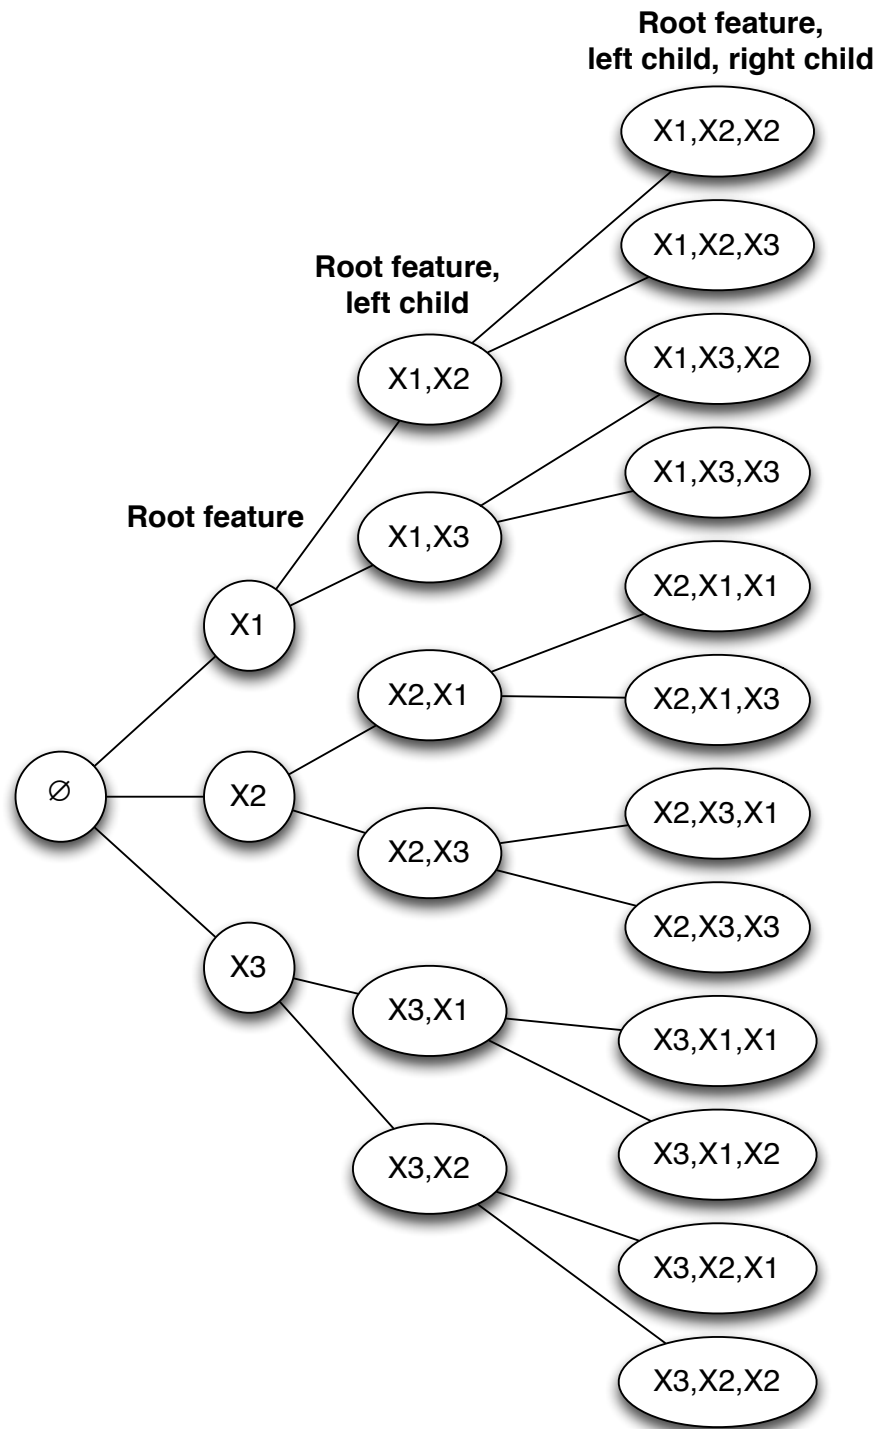
What is search space for decision-tree using data with three binary attributes  $X_1$ ,  $X_2$ ,  $X_3$



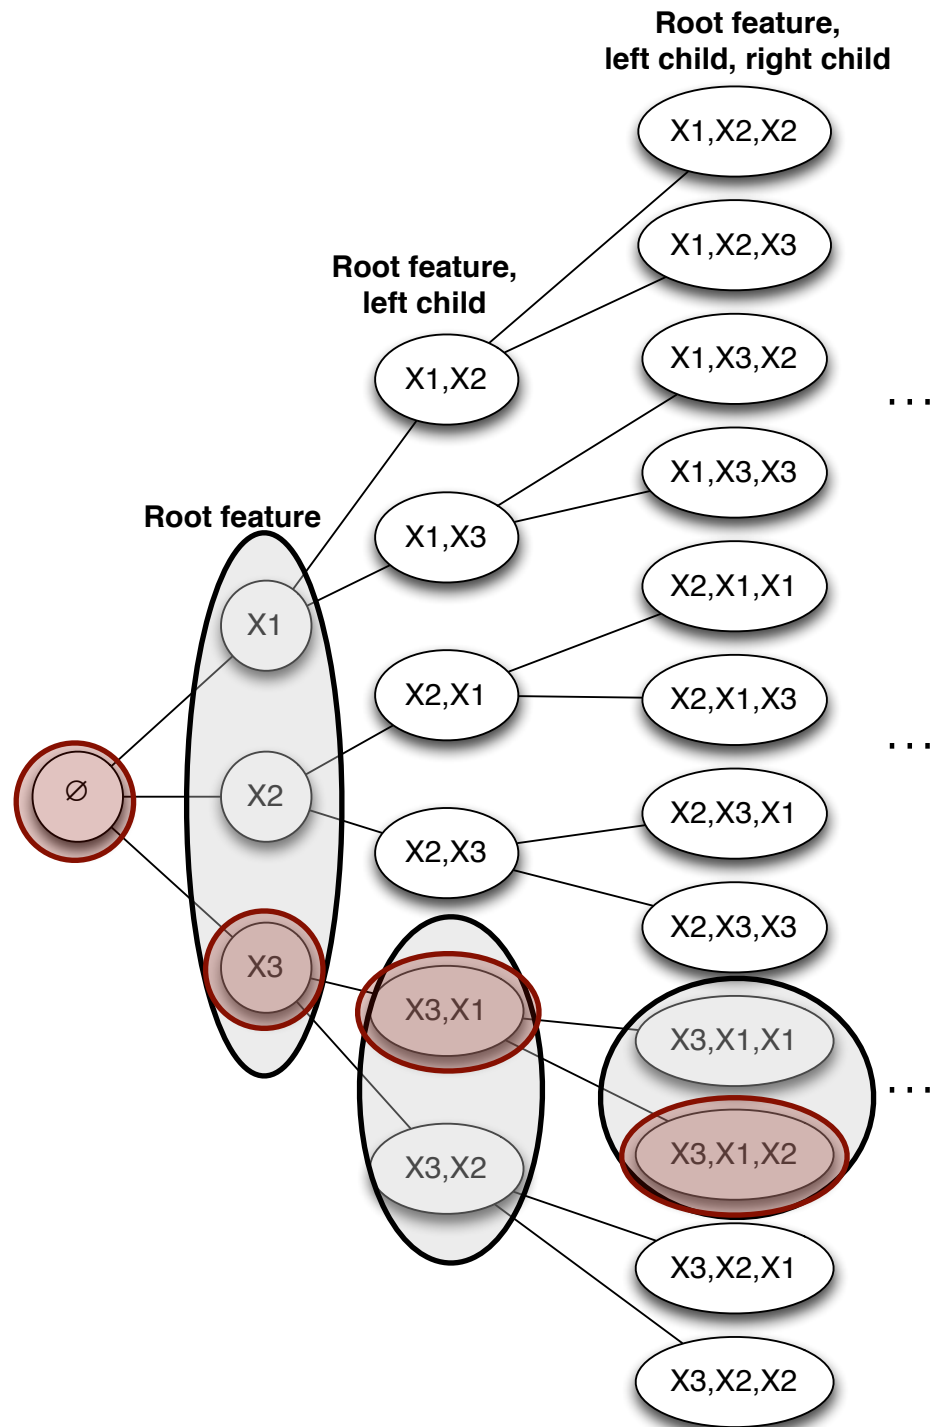
# Tree search algorithm

---

- Input:
  - Initial state? State space?
  - Set of actions?
  - How to choose next state?
  - Goal test?
- Output:
  - ?



Which states  
does greedy  
search  
consider?



# Questions to ask about search procedures

---

- Is the search **exhaustive**?
  - I.e., does it either *explicitly* or *implicitly* consider all models in the space?
- Is the search **optimal**?
  - I.e., is it *guaranteed* to return the model with the best score?
  - Global vs. local optimum?

Smooth optimization

# Optimization

---

- **Smooth** functions:
  - If a function is *smooth*, it is differentiable and the derivatives are continuous, then we can use gradient-based optimization
    - If function is *convex*, we can solve the minimization problem in closed form:  $\nabla S(\theta)$  using **convex optimization**
    - If function is smooth but non-linear, we can use iterative search over the surface of  $S$  to find a local minimum (e.g., hill-climbing)

# Convex optimization problems

---

$$\begin{array}{ll}\text{minimize} & f(x) \\ \text{subject to} & x \in C\end{array}$$

- Where  $f$  is a convex function (*score function*)  
 $C$  is a convex set (*constraints on model parameters or structure*)  
 $x$  is the optimization variable (*includes data and parameters*)
- For convex optimization problems, all locally optimal points are globally optimal
- Example algorithms: Quadratic programming (SVMs), least squares estimation, *maximum likelihood estimation*


Convex optimization



# Convex functions

---

- In graph of convex function, the line connecting two points must lie above the function



A function  $f$  is *convex* if:


$$f(\alpha x + (1 - \alpha)y) \leq \alpha f(x) + (1 - \alpha)f(y) \text{ for all } 0 \leq \alpha \leq 1$$

# Concave vs convex


---

- Maximizing a concave function is equivalent to minimizing a convex function

**concave**



**convex**



# Testing concavity (for maximization)

---

For a function  $f$  with parameters  $\Theta$

- Convexity conditions:

If  $f$  is continuous on  $\Theta$  and  $f''(\theta) \leq 0$  for all interior points  $\theta$

Then  $\mathbf{f}$  is a strictly concave function

- Sufficient conditions:

If  $f'(\theta) = 0$  then we say  $\theta$  is a *stationary point* of  $f$

If  $f'(\theta) = 0$  and  $f''(\theta) \leq 0$  then  $\theta$  is a *local maximum* of  $f$

If  $\mathbf{f}$  is a strictly concave function, any stationary point of  $\mathbf{f}$  is the unique global maximum of  $\mathbf{f}$

# Score function: Likelihood

---

- Let  $D = \{x(1), \dots, x(n)\}$
- Assume the data  $D$  are independently sampled from the same distribution:

$$p(X|\theta)$$

- The likelihood function represents the probability of the data as a function of the model parameters:

$$\begin{aligned} L(\theta|D) &= L(\theta|x(1), \dots, x(n)) \\ &= p(x(1), \dots, x(n)|\theta) \\ &= \prod_{i=1}^n p(x(i)|\theta) \end{aligned}$$

**If instances are independent,  
likelihood is product of probs**

# Maximum likelihood estimation

---

- Most widely used method of parameter estimation
- “Learn” the best parameters by finding the values of  $\theta$  that maximizes likelihood:

$$\hat{\theta}_{MLE} = \arg \max_{\theta} L(\theta)$$

- Often easier to work with loglikelihood:

$$\begin{aligned} l(\theta|D) &= \log L(\theta|D) \\ &= \log \prod_{i=1}^n p(x(i)|\theta) \\ &= \sum_{i=1}^n \log p(x(i)|\theta) \end{aligned}$$

# Maximum likelihood estimation

---

- Define likelihood, take derivative, set to 0, and solve
- Example
  - Toss a weighted coin 100 times, observe 30 heads
  - What is the MLE estimate for the  $p$  parameter of the Binomial distribution that generated the data?
  - First define likelihood

$$\begin{aligned} L(p|H=30, n=100) &= P(H=30|n=100, p) \\ &= \binom{100}{30} p^{30} (1-p)^{70} \end{aligned}$$

## Example (cont')

---

**Take derivative, set to 0, solve**

$$\begin{aligned} 0 &= \frac{d}{dp} \left( \binom{100}{30} p^{30} (1-p)^{70} \right) \\ &\propto 30p^{29} (1-p)^{70} - 70p^{30} (1-p)^{69} \\ &= p^{29} (1-p)^{69} [30(1-p) - 70p] \\ &= p^{29} (1-p)^{69} [30 - 100p] \\ p &= 0, 1, \frac{30}{100} \end{aligned}$$

# Gradient descent

---

- For some convex functions, we may be able to take the derivative, but it may be difficult to directly solve for parameter values
- Solution:
  - Start at some value of the parameters
  - Take derivative and use it to move the parameters in the direction of the solution
  - Repeat

## Gradient Descent Rule:

$$\underline{\mathbf{w}}_{\text{new}} = \underline{\mathbf{w}}_{\text{old}} - \eta \Delta(\underline{\mathbf{w}})$$

where

$\Delta(\underline{\mathbf{w}})$  is the gradient and

$\eta$  is the learning rate (small, positive)


Notes:

1. This moves us downhill in direction  $\Delta(\underline{\mathbf{w}})$  (steepest downhill direction)
2. How far we go is determined by the value of  $\eta$




# Illustration of gradient descent

---




# Illustration of gradient descent

---




# Illustration of gradient descent

---



# Illustration of gradient descent




Non-convex optimization

# Gradient descent

---


- Works on any objective function  $F(w)$ 
  - as long as we can evaluate the gradient  $\Delta(w)$
  - this can be very useful for minimizing complex functions  $F$
- Can be used in hill-climbing search to find local minima in smooth, but non-convex functions



- If function has multiple local minima, gradient descent goes to the closest local minimum:
  - solution: random restarts from multiple places in weight space

# Gradient ascent example

---



# Local search

---

- Iterative search that continually moves in direction of increasing value
  - Examples: gradient descent, hill climbing, coordinate descent
- Consider state-space landscape where:
  - Location=state; Elevation=**score function** for state

