

Data Mining & Machine Learning

CS37300
Purdue University


October 6, 2017

Extra Credit Competition Update

Public Leaderboard Private Leaderboard

This leaderboard is calculated with approximately 30% of the test data.
The final results will be based on the other 70%, so the final standings may be different.

 Raw Data  Refresh

#	Δ1w	Team Name	Kernel	Team Members	Score 	Entries	Last
1	new	Luke Skywalker	NBC		0.74585	2	2d
2	new	Bossk			0.67826	2	21m
3	new	Yoda			0.63820	1	9h
4	new	Revan			0.55281	2	11h
5	new	Boba Fett			0.51800	3	11h
		Bank_Sample_Submission.csv			0.49008		

How to beat Skywalker...

or

Classifiers beyond NBC and Decision Trees

So far...

A few weeks ago...
we reviewed Naive Bayes Classifier
and the Decision Tree...

We now embark on a quest
to find other classifiers

Classifiers for today

- Nearest neighbors
- Linear Regression
- Support vector machines
- Logistic Regression (1-layer neural network)

Classification Task

- ▶ Data representation:

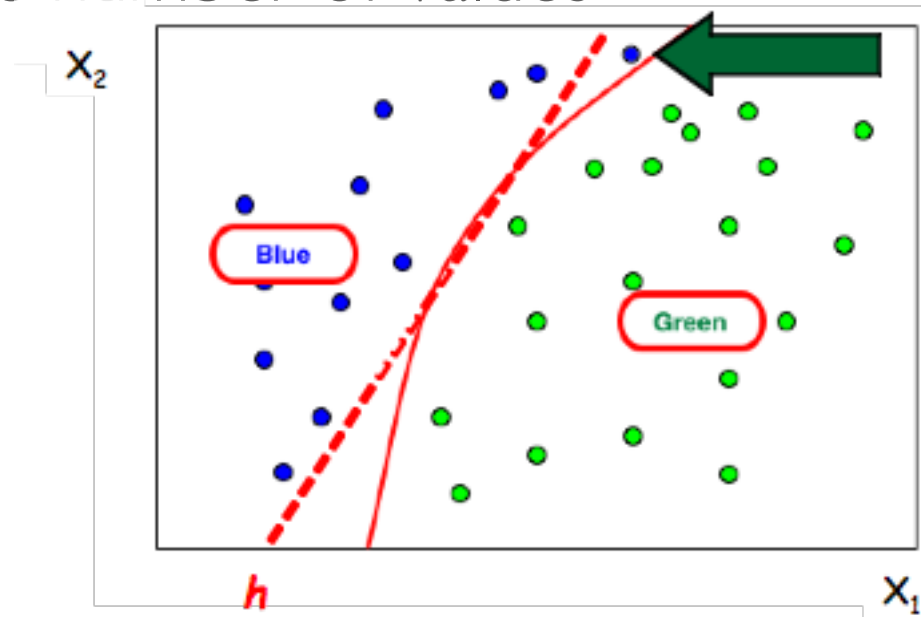
- Training set: Paired attribute vectors and class labels $\langle y(i), \mathbf{x}(i) \rangle$
or
 $n \times p$ tabular data with class label (y) and $p-1$ attributes (\mathbf{x})

- ▶ Task: estimate a predictive function $f(\mathbf{x}; \theta) = y$

- Assume that there is a function $y = f(\mathbf{x})$ that maps data instances (\mathbf{x}) to class labels (y)
- Construct a model that approximates the mapping
 - Classification: if y is categorical (e.g., {yes, no}, {dog, cat, elephant})
 - Regression: if y is real-valued (e.g., stock prices)

Binary classification

- ▶ In its simplest form, a classification model defines a decision boundary (h) and labels for each side of the boundary
- ▶ Input: $\mathbf{x} = \{x_1, x_2, \dots, x_n\}$ is a set of attributes, function f assigns a label y to input \mathbf{x} , where y is a discrete variable with a finite number of values



Nearest Neighbors

Nearest neighbor

- Instance-based method
- Learning
 - Stores training data and delays processing until a new instance must be classified
 - Assumes that all points are represented in p -dimensional space
- Prediction
 - **Nearest neighbors** are calculated using Euclidean distance
 - Classification is made based on class labels of neighbors

1NN

- Training set: $(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_n, y_n)$
where $\mathbf{x}_i = [x_{i1}, x_{i2}, \dots, x_{ip}]$ is a feature vector of p continuous attributes
and y_i is a discrete class label

- **1NN algorithm**

To predict a class label for new instance j :


Find the training instance point \mathbf{x}_i such that $d(\mathbf{x}_i, \mathbf{x}_j)$ is minimized

Let $f(\mathbf{x}_j) = y_i$


- Key idea: Find instances that are “similar” to the new instance and use their class labels to make prediction for the new instance
 - 1NN generalizes to kNN when more neighbors are considered

kNN model: decision boundaries

the data



NN classifier



kNN

- **kNN algorithm**

To predict a class label for new instance j :

Find the k nearest neighbors of j , i.e., those that minimize $d(\mathbf{x}_k, \mathbf{x}_j)$


Let $f(\mathbf{x}_j) = g(\mathbf{y}_k)$, e.g., majority label in \mathbf{y}_k

- *Algorithm choices*

- How many neighbors to consider (i.e., choice of k)?
... Usually a small value is used, e.g. $k < 10$
- What distance measure $d()$ to use?
... Euclidean L2 distance is often used
- What function $g()$ to combine the neighbors' labels into a prediction?
... Majority vote is often used

1NN decision boundary

- For each training example i , we can calculate its **Voronoi cell**, which corresponds to the space of points for which i is their nearest neighbor
- All points in such a Voronoi cell are labeled by the class of the training point, forming a Voronoi tessellation of the feature space



Nearest neighbor

- Strengths:
 - Simple model, easy to implement
 - Very efficient learning: $O(1)$
- Weaknesses:
 - Inefficient inference: time and space $O(n)$
 - Curse of dimensionality:
 - As number of features increase, you need an exponential increase in the size of the data to ensure that you have nearby examples for any given data point

k-NN learning

- Parameters of the model:
 - k (number of neighbors)
 - any parameters of distance measure (e.g., weights on features)
- **Model space**
 - Possible tessellations of the feature space
- **Search algorithm**
 - Implicit search: choice of k , d , and g uniquely define a tessellation
- **Score function**
 - Majority vote is minimizing misclassification rate

Least Squares Classifier

Motivation

- ▶ Given x features of a car (length, width, mpg, maximum speed,...)
- ▶ Classify cars into categories based on x

small car rentals ›



compacts
economy car rentals

medium car & SUV rentals ›



Coupes
Sedans
intermediate
SUV rentals

large car & SUV rentals ›



standard SUVs
premiums
luxury car rentals

fuel efficient & hybrid ›



Green car rentals

high occupancy car rentals ›



12-passenger vans
mini vans
premium SUV rentals

reservable models ›



Corvettes
Infinitis
BMW's & more

Least Squares Classifier

Two classes:

- ▶ \mathbf{x} is a real-valued vector (features)
- ▶ y is the car class

$$y_c = \begin{cases} 1 & , \text{ if car } c \text{ is "economy"} \\ -1 & , \text{ if car } c \text{ is "luxury"} \end{cases}$$


- ▶ Find linear discriminant weights \mathbf{w}

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$$

- ▶ Score function **least squares error**

$$\text{score}(\text{Test_Data_Cars}) = \sum_{c \in \text{Test_Data_Cars}} (y_c - y(x_c))^2$$



- ▶ Search function: **find \mathbf{w} , w_0 that minimize score**



Issues with Least Squares Classification

$$\text{score}(\text{Test_Data_Cars}) = \sum_{c \in \text{Test_Data_Cars}} (y_c - y(x_c))^2$$

Least Squares Solution




cares too much about well classified items

Neural networks

- Analogous to biological systems
- Massive parallelism is computationally efficient
- First learning algorithm in 1959 (Rosenblatt)
 - Perceptron learning rule
 - Provide target outputs with inputs for a single neuron
 - Incrementally update weights to learn to produce outputs

Neuron



$$f(x) = \sum_{i=1}^m w_i x_i + b$$

Activation function

$$y(x) = \sigma(f(x))$$

Single neuron (Logistic regression)

- ▶ **Model:** Single neuron is often used for two classes ($y=0, y=1$)

$$p(y = 1|x) = \sigma(f(x))$$

where

$$\sigma(a) = \frac{\exp(a)}{1 + \exp(a)} \quad \text{Logistic function}$$

$$f(x) = \sum_{i=1}^m w_i x_i + b$$

- ▶ **Score** function:

$$P[\text{Training Data} | \mathbf{w}, b] = \prod_{c \in \text{Training Data}} p(y(x_c) = y_c | x_c)$$


- ▶ **Search:** find w, b that minimize score

Support vector machines (SVMs)


Support vector machines

- Discriminative model
- General idea:
 - Find best boundary points (support vectors) and build classifier on top of them
- Linear and non-linear SVMs

Choosing hyperplanes to separate points




Among equivalent hyperplanes, choose one that maximizes “margin”



Linear SVMs

$$y = \text{sign} \left[\sum_{i=1}^m w_i x_i + b \right]$$

- Same functional form as perceptron
- Different learning procedure:
Search for hyperplane with largest margin
- Margin = $d_+ + d_-$
where d_+ is distance to closest positive example and d_- is distance to closest negative example



Constrained optimization for SVMs

$$Eq1 : x(j) \cdot w + b \geq +1 \text{ for } y(j) = +1$$

$$Eq2 : x(j) \cdot w + b \leq -1 \text{ for } y(j) = -1$$

Prediction
constraint

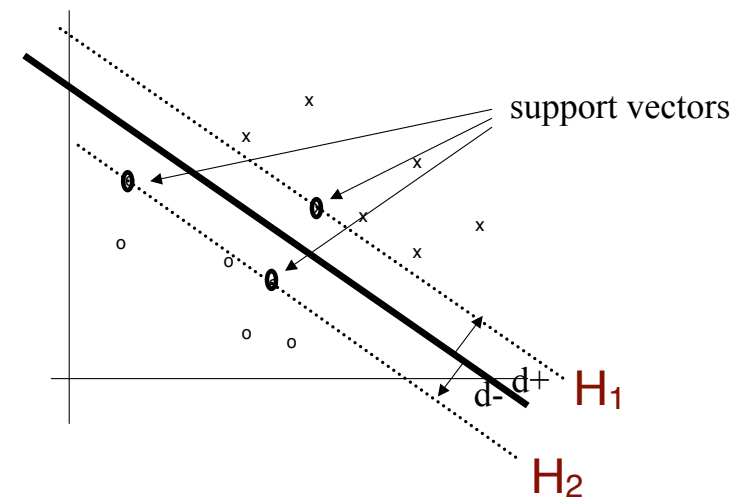
$$Eq3 : y(j)(x(j) \cdot w + b) - 1 \geq 0 \quad \forall y(j)$$

Hyperplane
boundaries

$$H_1 : x(j) \cdot w + b = +1$$

$$H_2 : x(j) \cdot w + b = -1$$

$$d_+ = d_- = \frac{1}{\|w\|} \qquad \text{margin} = \frac{2}{\|w\|}$$



- Can maximize margin by minimizing $\|w\|$ as it defines the hyperplanes

SVM optimization

- **Search:** Maximize margin by minimizing $0.5\|w\|^2$ subject to constraints on Eq3
 - Note: Maximizing $2/\|w\|$ is equivalent to minimizing $0.5\|w\|^2$



- Introduce Lagrange multipliers (α) for constraints into score function to minimize:

$$L_P = \frac{1}{2}\|w\|^2 - \sum_{i=1}^I \alpha_i y(i) [x(i) \cdot w + b] + \sum_{i=1}^I \alpha_i$$

- Minimize L_P with respect to w , b , and $\alpha_N \geq 0$
- Convex programming problem
 - Quadratic programming problem with parameters w , b , α

Constrained optimization

- Linear programming (LP) is a technique for the optimization of a linear objective function, subject to linear constraints on the variables
- Quadratic programming (QP) is a technique for the optimization of a quadratic function of several variables, subject to linear constraints on these variables



SVM components

- **Model space**
 - Set of weights \mathbf{w} and b (hyperplane boundary)
- **Search algorithm**
 - Quadratic programming to minimize L_p with constraints
- **Score function**
 - L_p : maximizes margin subject to constraint that all training data is correctly classified

Limitations of linear SVMs

- Linear classifiers cannot deal with:
 - Non-linear concepts
 - Noisy data
- Solutions:
 - Soft margin (e.g., allow mistakes in training data)
 - Network of simple linear classifiers (e.g., neural networks)
 - Map data into richer feature space (e.g., non-linear features) and then use linear classifier