

# Data Mining & Machine Learning

---

CS37300

Purdue University

November 1, 2017

# Model Search Tricks

Midsemester Course Evaluation (+1 % points)

Instructions on piazza


# Outline

---

- Regularization
  - L2 regularization
  - Dropout
- Other Tricks

# Neural Networks: Parameters x Training data sizes

- Consider a simple convolutional neural network for handwritten black/white digit classification:
  - The training data has only 50,000 images (and 50,000 labels).
  - Network has 1.2M parameters
  - The network likely will overfit the training data (not necessarily a problem, but could be)

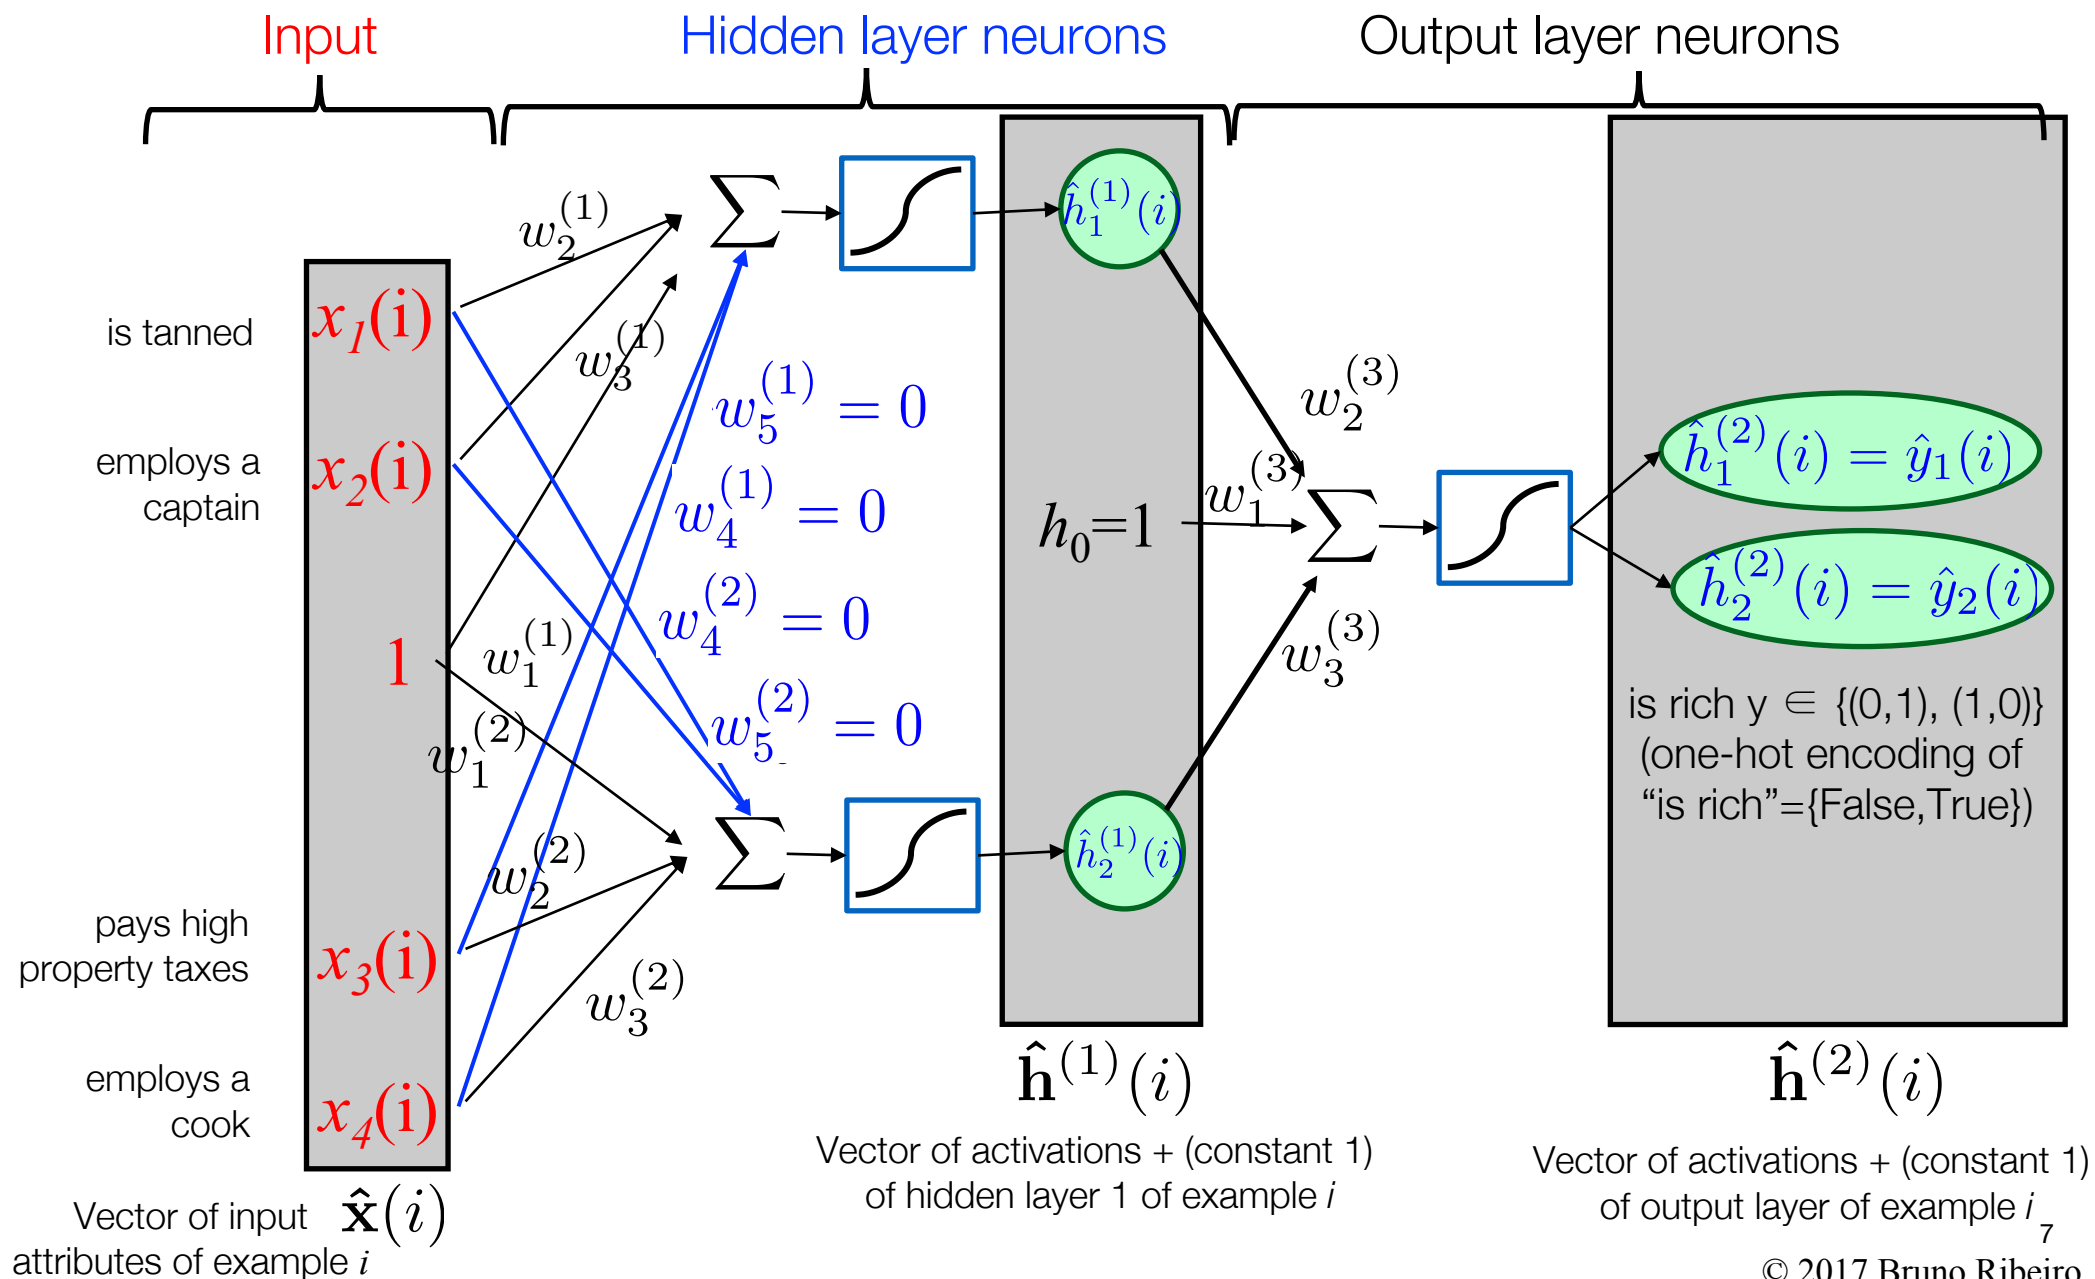


# Regularization

---

- There are many ways to avoid overfitting *too much*
  - Deep neural networks generally will overfit
- The idea behind regularization is to constraint either in the **model space** or the **model search**
  - In model space: E.g., forcing some parameters to be zero constrains the model
  - In model search: E.g., penalizing weights with large values constrains the search, and thus, the model

# Constraining the **Model Space** (Illustration)



# Constraining the **Model Search** (illustration)

---

- Constraining the model search:
  - In the previous example we pre-defined specific weights to be zero (constraining the model)
  - We could instead force 4 weights to be zero, but letting search algorithm decide which weights will be zero
  - This can improve the training accuracy of the model. Why?
    - If the the 4 pre-defined weights give better models, then the a good search algorithm will decide that these 4 weights give a better model
  - Asking for exactly 4 weights = zero is undesirable
    - Model should have some extra flexibility in deciding what is best
    - Answer: Penalties, penalize large weights (or the number of non-zero weights)



# Constraining Model Search in Practice

Let  $f(x; \mathbf{w})$  be the model the classifier that outputs the label of  $x$  given parameters  $\mathbf{w}$ . The training examples are  $\{x(i), y(i)\}_{i=1}^n$ , where  $x(i)$  are the attributes of observation  $i$  and  $y(i)$  its label. Suppose that  $S(y, y')$  is a score function between label  $y$  and predicted probability  $y'$ : better models have larger values (e.g., likelihood function). The model learning can be described as

$$\arg \max_{\mathbf{w}} \frac{1}{n} \sum_i S(y(i), f(x(i); \mathbf{w})) - \lambda \sum_j \sum_k (\mathbf{w}_j^{(k)})^2,$$

The parameter  $\lambda$  is known as the regularization strength. a.k.a.  $L_2$  norm  
penalizes large weights

Other alternative penalties include:

$$L_1 \text{ norm} = - \sum_j \sum_k |\mathbf{w}_j^{(k)}| \quad \text{Penalizes the absolute weight values}$$

$$L_0 \text{ norm} = - \sum_j \sum_k \mathbf{1}\{|\mathbf{w}_j^{(k)}| > 0\} \quad \text{Penalizes the number of non-zero weights}$$

# Regularization Strength

---

- Parameter  $\lambda$  is the regularization strength
  - How much we are penalizing the model?

$$\arg \max_{\mathbf{w}} \frac{1}{n} \sum_i S(y(i), f(x(i); \mathbf{w})) - \lambda \sum_j \sum_k (\mathbf{w}_j^{(k)})^2$$

Q: What happens if  $\lambda \rightarrow \infty$ ?

What happens if  $\lambda \rightarrow 0$ ?

A:

For  $\lambda \rightarrow \infty$ , the penalty for non-zero weights is too large, weights will be zero. The model is too constrained.

For  $\lambda \rightarrow 0$ , there is no penalty. The model is unconstrained.

Dropout (Model Space constraint)

# Averaging Many Models

---

- To win a machine learning competition (e.g. Netflix) you often need to combine many different types of strong models (complex models) and combine them to make predictions
- For instance, decision trees are not very powerful models, but
  - Averaging many decision trees works very well with some random model constraints. This is called random forest.
  - Random model constraints include only being able to split each decision tree node using a random subset of the attributes
    - This is a random constraint in the model space

# Two ways to average models

---

- Mixture: We combine the models by taking the arithmetic mean of their output probabilities (probability of a class).

Model A:    .3   .2   .5

Model B:    .1   .8   .1

---

Combined:   .2   .5   .3

- Product: Combine the models by taking the geometric mean of their probability.

Model A:    .3    .2    .5


Model B:    .1    .8    .1

---

Combined:  $\sqrt{.3}$   $\sqrt{.16}$   $\sqrt{.5}$  (need to renormalize to get a probability again)

# Dropout: An efficient way to average many large neural nets.

- Consider a neural net with one hidden layer.
- Each time we present a training example, we randomly omit each hidden unit with probability 0.5.
- So we are randomly sampling from  $2^H$  different architectures.
  - All architectures share weights.



Performed in training, validation, and testing

Hinton

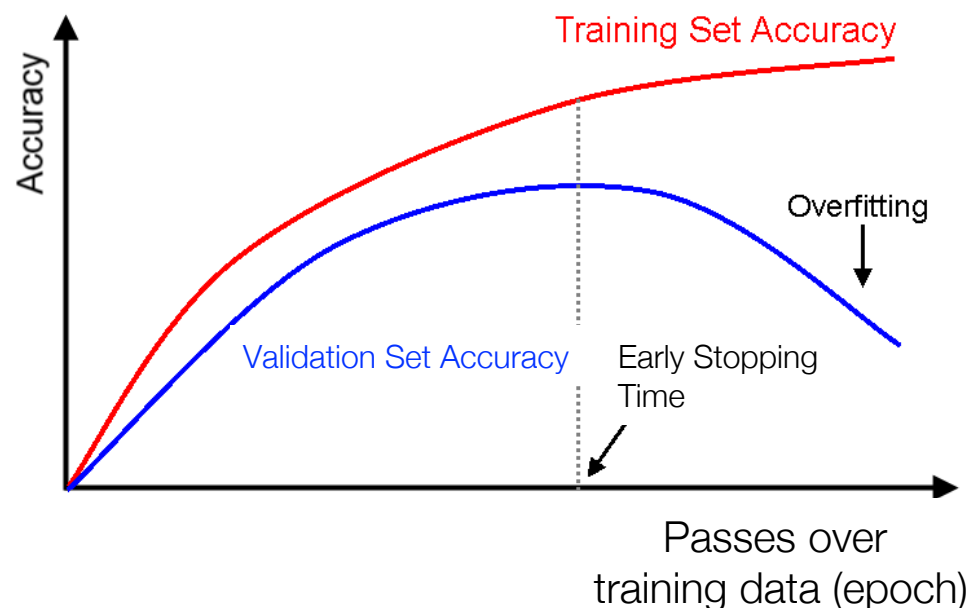
# Dropout as a form of model averaging

- We sample from  $2^H$  models. So only a few of the models ever get trained, and they only get one training example.
  - This is as extreme as bagging can get.
- The sharing of the weights means that every model is very strongly regularized.
  - It's a much better regularizer than L2 or L1 penalties that pull the weights towards zero.
- The neural network classifier output is the average over multiple different random dropouts

# Other tricks 1

---

- Early stopping:
  - Monitor the model accuracy in a separate validation dataset
  - As we are performing gradient ascent, if the model accuracy first increases but then starts to drop in the validation data, we stop training the model...
  - The output is the model parameter with the best accuracy in the validation data





# Other tricks 2

---

- Momentum
  - We will see it in HW 5